

# Week 4: Graph Data Structures

## Graph Definitions

### Graphs

2/106

Many applications require

- a collection of *items* (i.e. a set)
- *relationships*/connections between items

Examples:

- maps: items are cities, connections are roads
- web: items are pages, connections are hyperlinks

Collection types you're familiar with

- arrays, lists ... linear sequence of items (last week; COMP9021)

*Graphs* are more general ... allow arbitrary connections


### ... Graphs

3/106

A graph  $G = (V, E)$

- $V$  is a set of *vertices*
- $E$  is a set of *edges* (subset of  $V \times V$ )

Example:



$$V = \{v1, v2, v3, v4\}$$

$$E = \{e1, e2, e3, e4, e5\}$$

### ... Graphs

4/106

A real example: Australian road distances

Distance	Adelaide	Brisbane	Canberra	Darwin	Melbourne	Perth	Sydney
Adelaide	-	2055	-	3051	732	2716	-
Brisbane	2055	-	-	3429	1671	-	982


Canberra	-	-	-	-	658	-	309
Darwin	3051	3429	-	-	-	4049	-
Melbourne	732	1671	658	-	-	-	873
Perth	2716	-	-	4049	-	-	-
Sydney	-	982	309	-	873	-	-

Notes: vertices are cities, edges are distance between cities, symmetric

### ... Graphs

5/106

Alternative representation of above:



### ... Graphs

6/106

Questions we might ask about a graph:

- is there a way to get from item A to item B?
- what is the best way to get from A to B?
- which items are connected?

Graph algorithms are generally more complex than tree/list ones:

- no implicit order of items
- graphs may contain cycles
- concrete representation is less obvious
- algorithm complexity depends on connection complexity

## Properties of Graphs

7/106

Terminology:  $|V|$  and  $|E|$  (cardinality) normally written just as  $V$  and  $E$ .

A graph with  $V$  vertices has at most  $V(V-1)/2$  edges.

The ratio  $E:V$  can vary considerably.

- if  $E$  is closer to  $V^2$ , the graph is *dense*
- if  $E$  is closer to  $V$ , the graph is *sparse*
  - Example: web pages and hyperlinks

Knowing whether a graph is sparse or dense is important

- may affect choice of data structures to represent graph
- may affect choice of algorithms to process graph

## Exercise #1: Number of Edges

8/106

The edges in a graph represent pairs of connected vertices. A graph with  $V$  has  $V^2$  such pairs.

Consider  $V = \{1,2,3,4,5\}$  with all possible pairs:

$E = \{ (1,1), (1,2), (1,3), (1,4), (1,5), (2,1), (2,2), \dots, (4,5), (5,5) \}$

Why do we say that the maximum #edges is  $V(V-1)/2$ ?

... because

- $(v,w)$  and  $(w,v)$  denote the same edge (in an undirected graph)
- we do not consider loops  $(v,v)$

## Graph Terminology

10/106

For an edge  $e$  that connects vertices  $v$  and  $w$

- $v$  and  $w$  are *adjacent* (neighbours)
- $e$  is *incident* on both  $v$  and  $w$

*Degree* of a vertex  $v$

- number of edges incident on  $v$

Synonyms:

- vertex = node, edge = arc = link (Note: some people use arc for *directed* edges)

## ... Graph Terminology

11/106

*Path*: a sequence of vertices where

- each vertex has an edge to its predecessor

*Simple path*: a path where


- all vertices and edges are different

*Cycle*: a path

- that is simple except last vertex = first vertex

*Length* of path or cycle:

- #edges



## ... Graph Terminology

12/106

*Connected graph*

- there is a *path* from each vertex to every other vertex
- if a graph is not connected, it has  $\geq 2$  *connected components*

*Complete graph  $K_V$*

- there is an *edge* from each vertex to every other vertex
- in a complete graph,  $E = V(V-1)/2$



## ... Graph Terminology

13/106

*Tree*: connected (sub)graph with no cycles

*Spanning tree*: tree containing all vertices

*Clique*: complete subgraph

Consider the following single graph:



This graph has 26 vertices, 33 edges, and 4 connected components

Note: The entire graph has no spanning tree; what is shown in green is a spanning tree of the third connected component


A *spanning tree* of connected graph  $G = (V,E)$

- is a subgraph of  $G$  containing all of  $V$
- and is a single tree (connected, no cycles)

A *spanning forest* of non-connected graph  $G = (V,E)$

- is a subgraph of  $G$  containing all of  $V$
- and is a set of trees (not connected, no cycles),
  - with one tree for each *connected component*

Exercise #2: Graph Terminology



1. How many edges to remove to obtain a spanning tree?
2. How many different spanning trees?

1. 2
2.  $\frac{5 \cdot 4}{2} - 2 = 8$  spanning trees (no spanning tree if we remove  $\{e1,e2\}$  or  $\{e3,e4\}$ )


*Undirected graph*

- $edge(u,v) = edge(v,u)$ , no self-loops (i.e. no  $edge(v,v)$ )

*Directed graph*

- $edge(u,v) \neq edge(v,u)$ , can have self-loops (i.e.  $edge(v,v)$ )

Examples:



Other types of graphs ...

*Weighted graph*

- each edge has an associated value (weight)
- e.g. road map (weights on edges are distances between cities)

*Multi-graph*

- allow multiple edges between two vertices
- e.g. function call graph ( $f()$  calls  $g()$  in several places)


Graph Data Structures

Graph Representations

Defining graphs:

- need some way of identifying vertices
- could give diagram showing edges and vertices
- could give a list of edges

E.g. four representations of the same graph:



We will discuss three different graph data structures:

1. Array of edges
2. Adjacency matrix
3. Adjacency list

Array-of-edges Representation

Edges are represented as an array of Edge values (= pairs of vertices)

- space efficient representation
- adding and deleting edges is slightly complex
- undirected: order of vertices in an `Edge` doesn't matter
- directed: order of vertices in an `Edge` encodes direction



For simplicity, we always assume vertices to be numbered  $0 \dots V-1$

### ... Array-of-edges Representation

23/106

Graph initialisation

```
newGraph(V):
|   Input   number of nodes V
|   Output new empty graph
|
|   g.nV = V    // #vertices (numbered 0..V-1)
|   g.nE = 0    // #edges
|   allocate enough memory for g.edges[]
|   return g
```

How much is enough? ... No more than  $V(V-1)/2$  ... Much less in practice (sparse graph)

### ... Array-of-edges Representation

24/106

Edge insertion

```
insertEdge(g, (v,w)):
|   Input graph g, edge (v,w)  // assumption: (v,w) not in g
|
|   g.edges[g.nE]=(v,w)
|   g.nE=g.nE+1
```

### ... Array-of-edges Representation

25/106

Edge removal

```
removeEdge(g, (v,w)):
|   Input graph g, edge (v,w)  // assumption: (v,w) in g
|
|   i=0
|   while (v,w)≠g.edges[i] do
|       i=i+1
|   end while
```

```
|   g.edges[i]=g.edges[g.nE-1] // replace (v,w) by last edge in array
|   g.nE=g.nE-1
```

## Cost Analysis

26/106

Storage cost:  $O(E)$

Cost of operations:

- initialisation:  $O(1)$
- insert edge:  $O(1)$  (assuming edge array has space)
- find/delete edge:  $O(E)$  (need to find edge in edge array)

If array is full on insert

- allocate space for a bigger array, copy edges across  $\Rightarrow O(E)$

If we maintain edges in order

- use binary search to insert/find edge  $\Rightarrow O(\log E)$

### Exercise #3: Array-of-edges Representation

27/106

Assuming an array-of-edges representation ...

Write an algorithm to output all edges of the graph


```
show(g):
|   Input graph g
|
|   for all i=0 to g.nE-1 do
|       print g.edges[i]
|   end for
```

Time complexity:  $O(E)$

## Adjacency Matrix Representation


29/106

Edges represented by a  $V \times V$  matrix



Undirected graph

A	0	1	2	3
0	0	1	0	1
1	1	0	0	1
2	0	0	0	1
3	1	1	1	0



Directed graph

A	0	1	2	3
0	0	0	0	1
1	1	0	0	1
2	0	0	0	0
3	0	0	1	0

## ... Adjacency Matrix Representation

30/106

### Advantages

- easily implemented as 2-dimensional array
- can represent graphs, digraphs and weighted graphs
  - graphs: symmetric boolean matrix
  - digraphs: non-symmetric boolean matrix
  - weighted: non-symmetric matrix of weight values

### Disadvantages:

- if few edges (sparse)  $\Rightarrow$  memory-inefficient

## ... Adjacency Matrix Representation

31/106

### Graph initialisation

```
newGraph(V):
|   Input  number of nodes V
|   Output new empty graph
|
|   g.nV = V    // #vertices (numbered 0..V-1)
|   g.nE = 0    // #edges
|   allocate memory for g.edges[][]
|   for all i,j=0..V-1 do
|       g.edges[i][j]=0    // false
|   end for
|   return g
```

## ... Adjacency Matrix Representation

32/106

### Edge insertion

```
insertEdge(g, (v,w)):
```

```
|   Input graph g, edge (v,w)
```

```
|   if g.edges[v][w]=0 then    // (v,w) not in graph
|       g.edges[v][w]=1        // set to true
|       g.edges[w][v]=1
|       g.nE=g.nE+1
|   end if
```

## ... Adjacency Matrix Representation

33/106

### Edge removal

```
removeEdge(g, (v,w)):
```

```
|   Input graph g, edge (v,w)
```

```
|   if g.edges[v][w]≠0 then    // (v,w) in graph
|       g.edges[v][w]=0        // set to false
|       g.edges[w][v]=0
|       g.nE=g.nE-1
|   end if
```

## Exercise #4: Show Graph

34/106

Assuming an adjacency matrix representation ...

Write an algorithm to output all edges of the graph (no duplicates!)

## ... Adjacency Matrix Representation

35/106

```
show(g):
|   Input graph g
|
|   for all i=0 to g.nV-2 do
|       for all j=i+1 to g.nV-1 do
|           if g.edges[i][j] then
|               print i-"-j"
|           end if
|       end for
|   end for
```

Time complexity:  $O(V^2)$

## Exercise #5:

36/106

Analyse storage cost and time complexity of adjacency matrix representation

Storage cost:  $O(V^2)$

If the graph is sparse, most storage is wasted.


#### Cost of operations:

- initialisation:  $O(V^2)$  (initialise  $V \times V$  matrix)
- insert edge:  $O(1)$  (set two cells in matrix)
- delete edge:  $O(1)$  (unset two cells in matrix)

### ... Adjacency Matrix Representation

38/106

A storage optimisation: store only top-right part of matrix.



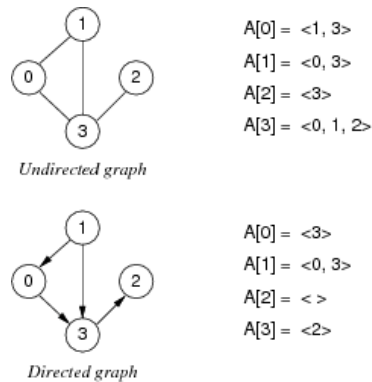
New storage cost:  $V-1$  int ptrs +  $V(V+1)/2$  ints (but still  $O(V^2)$ )

Requires us to always use edges  $(v,w)$  such that  $v < w$ .

### Adjacency List Representation

39/106

For each vertex, store linked list of adjacent vertices:



### ... Adjacency List Representation

40/106

Advantages

- relatively easy to implement in languages like C
- can represent graphs and digraphs

- memory efficient if  $E:V$  relatively small

Disadvantages:

- one graph has many possible representations (unless lists are ordered by same criterion e.g. ascending)

### ... Adjacency List Representation

41/106

Graph initialisation

```
newGraph(V):  
    Input number of nodes V  
    Output new empty graph  
  
    g.nV = V // #vertices (numbered 0..V-1)  
    g.nE = 0 // #edges  
    allocate memory for g.edges[]  
    for all i=0..V-1 do  
        g.edges[i]=NULL // empty list  
    end for  
    return g
```

### ... Adjacency List Representation

42/106

Edge insertion:

```
insertEdge(g, (v,w)):  
    Input graph g, edge (v,w)  
  
    insertLL(g.edges[v], w)  
    insertLL(g.edges[w], v)  
    g.nE=g.nE+1
```

### ... Adjacency List Representation

43/106

Edge removal:

```
removeEdge(g, (v,w)):  
    Input graph g, edge (v,w)  
  
    deleteLL(g.edges[v], w)  
    deleteLL(g.edges[w], v)  
    g.nE=g.nE-1
```

### Exercise #6:

44/106

Analyse storage cost and time complexity of adjacency list representation

Storage cost:  $O(V+E)$  ( $V$  list pointers, total of  $2 \cdot E$  list elements)

Cost of operations:

- initialisation:  $O(V)$  (initialise  $V$  lists)
- insert edge:  $O(I)$  (insert one vertex into list)
  - if you don't check for duplicates
- find/delete edge:  $O(V)$  (need to find vertex in list)

If vertex lists are sorted

- insert requires search of list  $\Rightarrow O(V)$
- delete always requires a search, regardless of list order

## Comparison of Graph Representations

46/106

	array of edges	adjacency matrix	adjacency list
space usage	$E$	$V^2$	$V+E$
initialise	$I$	$V^2$	$V$
insert edge	$I$	$I$	$I$
find/delete edge	$E$	$I$	$V$

Other operations:

	array of edges	adjacency matrix	adjacency list
disconnected(v)?	$E$	$V$	$I$
isPath(x,y)?	$E \cdot \log V$	$V^2$	$V+E$
copy graph	$E$	$V^2$	$E$
destroy graph	$I$	$V$	$E$

## Graph Abstract Data Type

### Graph ADT

48/106

Data:

- set of edges, set of vertices

Operations:

- building: create graph, add edge

- deleting: remove edge, drop whole graph
- scanning: check if graph contains a given edge

Things to note:

- set of vertices is fixed when graph initialised
- we treat vertices as ints, but could be arbitrary Items

### ... Graph ADT

49/106

Graph ADT interface **graph.h**

```
// graph representation is hidden
typedef struct GraphRep *Graph;

// vertices denoted by integers 0..N-1
typedef int Vertex;

// edges are pairs of vertices (end-points)
typedef struct Edge { Vertex v; Vertex w; } Edge;

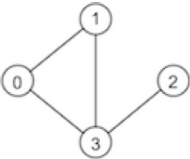
// operations on graphs
Graph newGraph(int V); // new graph with V vertices
void insertEdge(Graph, Edge);
void removeEdge(Graph, Edge);
bool adjacent(Graph, Vertex, Vertex); /* is there an edge
                                         between two vertices */
void freeGraph(Graph);
```

### Exercise #7: Graph ADT Client

50/106

Write a program that uses the graph ADT to

- build the graph depicted below
- print all the nodes that are incident to vertex 1 in ascending order



```
#include <stdio.h>
#include "Graph.h"

#define NODES 4
#define NODE_OF_INTEREST 1

int main(void) {
```

```

Graph g = newGraph(NODES);

Edge e;
e.v = 0; e.w = 1; insertEdge(g,e);
e.v = 0; e.w = 3; insertEdge(g,e);
e.v = 1; e.w = 3; insertEdge(g,e);
e.v = 3; e.w = 2; insertEdge(g,e);

int v;
for (v = 0; v < NODES; v++) {
    if (adjacent(g, v, NODE_OF_INTEREST))
        printf("%d\n", v);
}

freeGraph(g);
return 0;
}

```

## Graph ADT (Array of Edges)

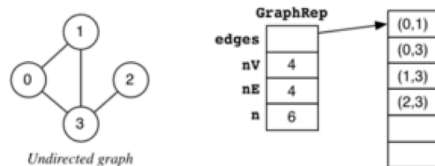
52/106

Implementation of GraphRep (array-of-edges representation)

```

typedef struct GraphRep {
    Edge *edges; // array of edges
    int nV;      // #vertices (numbered 0..nV-1)
    int nE;      // #edges
    int n;       // size of edge array
} GraphRep;

```



## Graph ADT (Adjacency Matrix)

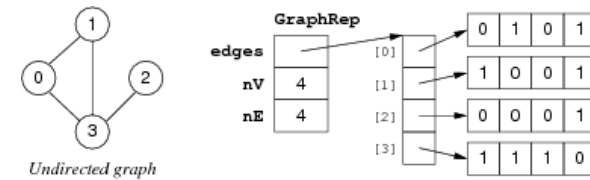
53/106

Implementation of GraphRep (adjacency-matrix representation)

```

typedef struct GraphRep {
    int **edges; // adjacency matrix
    int nV;      // #vertices
    int nE;      // #edges
} GraphRep;

```



## ... Graph ADT (Adjacency Matrix)

54/106

Implementation of graph initialisation (adjacency-matrix representation)

```

Graph newGraph(int V) {
    assert(V >= 0);
    int i;

    Graph g = malloc(sizeof(GraphRep));    assert(g != NULL);
    g->nV = V; g->nE = 0;

    // allocate memory for each row
    g->edges = malloc(V * sizeof(int *));    assert(g->edges != NULL);
    // allocate memory for each column and initialise with 0
    for (i = 0; i < V; i++) {
        g->edges[i] = calloc(V, sizeof(int)); assert(g->edges[i] != NULL);
    }
    return g;
}

```

standard library function **calloc(size\_t nelems, size\_t nbytes)**

- allocates a memory block of size `nelems*nbytes`
- and sets all bytes in that block to *zero*

## ... Graph ADT (Adjacency Matrix)

55/106

Implementation of edge insertion/removal (adjacency-matrix representation)

```

// check if vertex is valid in a graph
bool validV(Graph g, Vertex v) {
    return (g != NULL && v >= 0 && v < g->nV);
}

void insertEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));

    if (!g->edges[e.v][e.w]) { // edge e not in graph
        g->edges[e.v][e.w] = 1;
        g->edges[e.w][e.v] = 1;
        g->nE++;
    }
}

void removeEdge(Graph g, Edge e) {
    assert(g != NULL && validV(g,e.v) && validV(g,e.w));
}

```



```

    if (g->edges[e.v][e.w]) { // edge e in graph
        g->edges[e.v][e.w] = 0;
        g->edges[e.w][e.v] = 0;
        g->nE--;
    }
}

```

## Exercise #8: Checking Neighbours (i)

56/106

Assuming an adjacency-matrix representation ...

Implement a function to check whether two vertices are directly connected by an edge

```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```

```

bool adjacent(Graph g, Vertex x, Vertex y) {
    assert(g != NULL && validV(g,x) && validV(g,y));

    return (g->edges[x][y] != 0);
}

```

## Graph ADT (Adjacency List)

58/106

Implementation of GraphRep (adjacency-list representation)

```

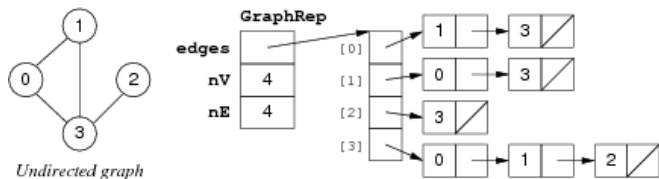
typedef struct GraphRep {
    Node **edges; // array of lists
    int    nV;    // #vertices
    int    nE;    // #edges
} GraphRep;

```

```

typedef struct Node {
    Vertex    v;
    struct Node *next;
} Node;

```



## Exercise #9: Checking Neighbours (ii)

59/106

Assuming an adjacency list representation ...

Implement a function to check whether two vertices are directly connected by an edge

```
bool adjacent(Graph g, Vertex x, Vertex y) { ... }
```

```

bool adjacent(Graph g, Vertex x, Vertex y) {
    assert(g != NULL && validV(g,x));

    return inLL(g->edges[x], y);
}

```

inLL() checks if linked list contains an element

## Problems on Graphs

## Problems on Graphs

62/106

What kind of problems do we want to solve on/via graphs?

- is the graph fully-connected?
- can we remove an edge and keep it fully-connected?
- is one vertex reachable starting from some other vertex?
- which vertices are reachable from  $v$ ? (transitive closure)
- is there a cycle that passes through all vertices? (circuit)
- what is the cheapest cost path from  $v$  to  $w$ ?
- is there a tree that links all vertices? (spanning tree)
- what is the minimum spanning tree?
- what is the maximal flow through a graph?
- ...
- can a graph be drawn in a plane with no crossing edges? (planar graphs)
- are two graphs "equivalent"? (isomorphism)

## Graph Algorithms

63/106

In this course we examine algorithms for

- graph traversal (simple graphs)
- reachability (directed graphs)
- minimum spanning trees (weighted graphs)
- shortest path (weighted graphs)
- maximum flow (weighted graphs)

## Graph Traversal

## Finding a Path

65/106

Questions on paths:

- is there a path between two given vertices (*src*,*dest*)?
- what is the sequence of vertices from *src* to *dest*?

Approach to solving problem:

- examine vertices adjacent to *src*
- if any of them is *dest*, then done
- otherwise try vertices two edges from *src*
- repeat looking further and further from *src*


Two strategies for graph traversal/search: *depth-first*, *breadth-first*

- DFS follows one path to completion before considering others
- BFS "fans-out" from the starting vertex ("spreading" subgraph)

## ... Finding a Path

66/106

Comparison of BFS/DFS search for checking if there is a path from *a* to *h* ...



Both approaches ignore some edges by remembering previously visited vertices.

## Depth-first Search

67/106

Depth-first search can be described recursively as

```
depthFirst(G,v):
1. mark v as visited
2. for each (v,w) ∈ edges(G) do
   if w has not been visited then
     depthFirst(w)
```

The recursion induces *backtracking*

## ... Depth-first Search

68/106

Recursive DFS path checking

```
hasPath(G,src,dest):
| Input graph G, vertices src,dest
```

```
Output true if there is a path from src to dest in G,
        false otherwise
```


```
mark all vertices in G as unvisited
return dfsPathCheck(G,src,dest)
```

```
dfsPathCheck(G,v,dest):
| mark v as visited
| if v=dest then           // found dest
|   return true
| else
|   for all (v,w) ∈ edges(G) do
|   | if w has not been visited then
|   |   return dfsPathCheck(G,w,dest) // found path via w to dest
|   | end if
|   end for
| end if
| return false           // no path from v to dest
```

## Exercise #10: Depth-first Traversal (i)

69/106

Trace the execution of `dfsPathCheck(G,0,5)` on:



Consider neighbours in ascending order

Answer:

0 - 1 - 2 - 3 - 4 - 5

## ... Depth-first Search

71/106

Cost analysis:

- all vertices marked as unvisited, each vertex visited at most once  $\Rightarrow$  cost =  $O(V)$
- visit all edges incident on visited vertices  $\Rightarrow$  cost =  $O(E)$ 
  - assuming an adjacency list representation

*Time complexity of DFS:*  $O(V+E)$  (adjacency list representation)

- the larger of  $V,E$  determines the complexity

... Depth-first Search

72/106

Note how different graph data structures affect cost:

- array-of-edges representation
  - visit all edges incident on visited vertices  $\Rightarrow$  cost =  $O(V \cdot E)$
  - cost of DFS:  $O(V \cdot E)$
- adjacency-matrix representation
  - visit all edges incident on visited vertices  $\Rightarrow$  cost =  $O(V^2)$
  - cost of DFS:  $O(V^2)$

For *dense graphs* ...  $E \cong V^2 \Rightarrow O(V+E) = O(V^2)$   
For *sparse graphs* ...  $E \cong V \Rightarrow O(V+E) = O(E)$

... Depth-first Search

73/106

Knowing whether a path exists can be useful

Knowing what the path is even more useful  
 $\Rightarrow$  record the previously visited node as we search through the graph (so that we can then trace path through graph)

Make use of global variable:

- `visited[ ]` ... array to store previously visited node, for each node being visited

... Depth-first Search

74/106

`visited[ ]` // store previously visited node, for each vertex  $0..nV-1$

```
findPath(G,src,dest):
    Input graph G, vertices src,dest

    for all vertices v∈G do
        visited[v]=-1
    end for
    visited[src]=src                // starting node of the path
    if dfsPathCheck(G,src,dest) then // show path in dest..src order
        v=dest
        while v≠src do
            print v "-"
            v=visited[v]
        end while
        print src
    end if


dfsPathCheck(G,v,dest):
    if v=dest then                // found edge from v to dest
        return true
    else
```

```
        for all (v,w)∈edges(G) do
            if visited[w]=-1 then
                visited[w]=v
                if dfsPathCheck(G,w,dest) then
                    return true    // found path via w to dest
                end if
            end if
        end for
    end if
    return false                  // no path from v to dest
```

Exercise #11: Depth-first Traversal (ii)

75/106

Show the DFS order in which we visit vertices in this graph when searching for a path from 0 to 6:



Consider neighbours in ascending order

0	0	3	5	3	1	5	4	7	8
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

Path: 6-5-1-0

... Depth-first Search

77/106

DFS can also be described non-recursively (via a *stack*):

```
hasPath(G,src,dest):
    Input graph G, vertices src,dest
    Output true if there is a path from src to dest in G,
           false otherwise

    mark all vertices in G as unvisited
    push src onto new stack s
    found=false
    while not found and s is not empty do
        pop v from s
        mark v as visited
```

```

if v=dest then
  found=true
else
  for each (v,w)∈edges(G) such that w has not been visited
    push w onto s
  end for
end if
end while
return found

```


Uses standard stack operations (push, pop, check if empty)

Time complexity is the same:  $O(V+E)$  (each vertex added to stack once, each element in vertex's adjacency list visited once)

## Exercise #12: Depth-first Traversal (iii)

78/106

Show how the stack evolves when executing `findPathDFS(g,0,5)` on:



Push neighbours in *descending* order ... so they get popped in ascending order

				4	5											
			3	5	5	5										
	1	2	4	4	4	4										
	4	4	4	4	4	4										
(empty)	→	0	→	5	→	5	→	5	→	5	→	5	→	5	→	5

## Breadth-first Search

80/106

Basic approach to breadth-first search (BFS):

- visit and mark current vertex
- visit all neighbours of current vertex
- then consider neighbours of neighbours

Notes:

- tricky to describe recursively
- a minor variation on non-recursive DFS search works  
⇒ switch the *stack* for a *queue*

## ... Breadth-first Search

81/106

BFS algorithm (records visiting order, marks vertices as visited when put *on* queue):

`visited[]` // array of visiting orders, indexed by vertex `0..nV-1`

```

findPathBFS(G,src,dest):
  Input graph G, vertices src,dest

  for all vertices v∈G do
    visited[v]=-1
  end for
  enqueue src into new queue q
  visited[src]=src
  found=false
  while not found and q is not empty do
    dequeue v from q
    if v=dest then
      found=true
    else
      for each (v,w)∈edges(G) such that visited[w]=-1 do
        enqueue w into q
        visited[w]=v
      end for
    end if
  end while
  if found then
    display path in dest..src order
  end if


```

Uses standard queue operations (enqueue, dequeue, check if empty)

## Exercise #13: Breadth-first Traversal

82/106

Show the BFS order in which we visit vertices in this graph when searching for a path from 0 to 6:



Consider neighbours in ascending order

0	0	0	2	5	0	5	5	3	-1
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]	[8]	[9]

Path: 6-5-0

... Breadth-first Search

84/106

Time complexity of BFS:  $O(V+E)$  (adjacency list representation, same as DFS)

BFS finds a "shortest" path

- based on minimum # edges between *src* and *dest*.
- stops with first-found path, if there are multiple ones

In many applications, edges are weighted and we want path

- based on minimum sum-of-weights along path *src* .. *dest*


We discuss weighted/directed graphs later.

Other DFS Examples

85/106

Other problems to solve via DFS graph search

- checking for the existence of a cycle
- determining which connected component each vertex is in



Exercise #14: Buggy Cycle Check

86/106

A graph has a cycle if

- it has a path of length > 1
- with start vertex *src* = end vertex *dest*
- and without using any edge more than once

We are not required to give the path, just indicate its presence.

The following DFS cycle check has two bugs. Find them.

```
hasCycle(G):
|   Input  graph G
|   Output true if G has a cycle, false otherwise
|
|   choose any vertex v∈G
|   return dfsCycleCheck(G,v)
|
dfsCycleCheck(G,v):
|   mark v as visited
|   for each (v,w)∈edges(G) do
|       if w has been visited then    // found cycle
|           return true
|       else if dfsCycleCheck(G,w) then
|           return true
|   end for
|   return false                      // no cycle at v
```

1. Only one connected component is checked.
2. The loop

```
for each (v,w)∈edges(G) do
```

should exclude the neighbour of v from which you just came, so as to prevent a single edge *w-v* from being classified as a cycle.

Computing Connected Components

88/106

Problems:

- how many connected subgraphs are there?
- are two vertices in the same connected subgraph?

Both of the above can be solved if we can

- build an array, one element for each vertex V
- indicating which connected component V is in
- componentOf[ ] ... array [0..nV-1] of component IDs

... Computing Connected Components

89/106

Algorithm to assign vertices to connected components:

```
components(G):
|   Input  graph G
|
|   for all vertices v∈G do
|       componentOf[v]=--1
```

```
end for
compID=0
for all vertices v∈G do
  if componentOf[v]=-1 then
    dfsComponents(G,v,compID)
    compID=compID+1
  end if
end for
```


```
dfsComponents(G,v,id):
  componentOf[v]=id
  for all vertices w adjacent to v do
    if componentOf[w]=-1 then
      dfsComponents(G,w,id)
    end if
  end for
```

Exercise #15: Connected components

90/106

Trace the execution of the algorithm

- 1. on the graph shown below
- 2. on the same graph but with the dotted edges added



Consider neighbours in ascending order

1.

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
-1	-1	-1	-1	-1	-1	-1	-1
0	-1	-1	-1	-1	-1	-1	-1
0	-1	0	-1	-1	-1	-1	-1
0	0	0	-1	-1	-1	-1	-1
0	0	0	1	-1	-1	-1	-1
...							
0	0	0	1	1	2	2	2

2.

[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]
-1	-1	-1	-1	-1	-1	-1	-1
0	-1	-1	-1	-1	-1	-1	-1
0	0	-1	-1	-1	-1	-1	-1
0	0	0	-1	-1	-1	-1	-1
...							
0	0	0	0	0	1	1	1

Hamiltonian and Euler Paths

Hamiltonian Path and Circuit

93/106

Hamiltonian path problem:

- find a path connecting two vertices  $v,w$  in graph  $G$
- such that the path includes each *vertex* exactly once

If  $v = w$ , then we have a *Hamiltonian circuit*

Simple to state, but difficult to solve (*NP*-complete)

Many real-world applications require you to visit all vertices of a graph:


- Travelling salesman
- Bus routes
- ...

Problem named after Irish mathematician, physicist and astronomer Sir William Rowan Hamilton (1805 — 1865)

... Hamiltonian Path and Circuit

94/106

Graph and two possible Hamiltonian paths:



... Hamiltonian Path and Circuit

95/106

Approach:

- generate all possible simple paths (using e.g. DFS)

- keep a counter of vertices visited in current path
- stop when find a path containing  $V$  vertices

Can be expressed via a recursive DFS algorithm

- similar to simple path finding approach, except
  - keeps track of path length; succeeds if length =  $v$
  - resets "visited" marker after unsuccessful path

## ... Hamiltonian Path and Circuit

96/106

Algorithm for finding Hamiltonian path:

visited[] // array [0..nV-1] to keep track of visited vertices

```


hasHamiltonianPath(G,src,dest):
|   for all vertices v∈G do
|       visited[v]=false
|   end for
|   return hamiltonR(G,src,dest,#vertices(G)-1)

hamiltonR(G,v,dest,d):
|   Input G      graph
|       v      current vertex considered
|       dest destination vertex
|       d      distance "remaining" until path found
|
|   if v=dest then
|       if d=0 then return true else return false
|   else
|       mark v as visited
|       for each unvisited neighbour w of v in G do
|           if hamiltonR(G,w,dest,d-1) then
|               return true
|           end if
|       end for
|   end if
|   mark v as unvisited           // reset visited mark
|   return false
  
```

## Exercise #16: Hamiltonian Path

97/106

Trace the execution of the algorithm when searching for a Hamiltonian path from 1 to 6:



Consider neighbours in ascending order

1-0-2-3-4-5-6	d≠0
1-0-2-3-4-5-7-8-9	no unvisited neighbour
1-0-2-3-4-5-7-9-8	no unvisited neighbour
1-0-2-3-4-7-5-6	d≠0
1-0-2-3-4-7-8-9	no unvisited neighbour
1-0-2-3-4-7-9-8	no unvisited neighbour
1-0-2-3-4-8-7-5-6	d≠0
1-0-2-3-4-8-7-9	no unvisited neighbour
1-0-2-3-4-8-9-7-5-6	✓


Repeat on your own with src=0 and dest=6

## ... Hamiltonian Path and Circuit

99/106

Analysis: worst case requires  $(V-1)!$  paths to be examined

Consider a graph with isolated vertex and the rest fully-connected



Checking hasHamiltonianPath( $g, x, 0$ ) for any  $x$

- requires us to consider every possible path
- e.g 1-2-3-4, 1-2-4-3, 1-3-2-4, 1-3-4-2, 1-4-2-3, ...
- starting from any  $x$ , there are  $3!$  paths  $\Rightarrow 4!$  total paths
- there is no path of length 5 in these  $(V-1)!$  possibilities

There is no known simpler algorithm for this task  $\Rightarrow$  NP-hard.

Note, however, that the above case could be solved in constant time if we had a fast check for 0 and  $x$  being in the same connected component


## Euler Path and Circuit

100/106


Euler path problem:

- find a path connecting two vertices  $v, w$  in graph  $G$
- such that the path includes each *edge* exactly once  
(note: the path does not have to be simple  $\Rightarrow$  can visit vertices more than once)

If  $v = w$ , then we have an *Euler circuit*



Euler Path: 4-3-1-5-2-1



Euler Circuit: 1-2-5-4-3-1

Many real-world applications require you to visit all edges of a graph:

- Postman
- Garbage pickup
- ...

Problem named after Swiss mathematician, physicist, astronomer, logician and engineer Leonhard Euler (1707 — 1783)

## ... Euler Path and Circuit

101/106

One possible "brute-force" approach:

- check for each path if it's an Euler path
- would result in factorial time performance

Can develop a better algorithm by exploiting:



*Theorem.* A graph has an Euler circuit if and only if it is connected and all vertices have even degree

*Theorem.* A graph has a non-circuitous Euler path if and only if it is connected and exactly two vertices have odd degree

## Exercise #17: Euler Paths and Circuits

102/106

Which of these two graphs have an Euler path? an Euler circuit?



No Euler circuit

Only the second graph has an Euler path, e.g. 2-0-1-3-2-4-5

## ... Euler Path and Circuit

104/106

Assume the existence of  $\text{degree}(g, v)$  (degree of a vertex, cf. assessment question 1 this week)

Algorithm to check whether a graph has an Euler path:

```
hasEulerPath(G,src,dest):
    Input  graph G, vertices src,dest
    Output true if G has Euler path from src to dest
           false otherwise

    if src≠dest then           // non-circuitous path
        if degree(G,src) or degree(G,dest) is even then
            return false
        end if
    else if degree(G,src) is odd then // circuit
        return false
    end if
    for all vertices v∈G do
        if v≠src and v≠dest and degree(G,v) is odd then
            return false
        end if
    end for
    return true
```

## ... Euler Path and Circuit

105/106

Analysis of hasEulerPath algorithm:

- assume that connectivity is already checked
- assume that degree is available via  $O(1)$  lookup
- single loop over all vertices  $\Rightarrow O(V)$

If degree requires iteration over vertices

- cost to compute degree of a single vertex is  $O(V)$
- overall cost is  $O(V^2)$

$\Rightarrow$  problem tractable, even for large graphs (unlike Hamiltonian path problem)

For the keen, a linear-time (in the number of edges,  $E$ ) algorithm to compute an Euler path is described in [Sedgewick] Ch.17.7.

## Summary

106/106

- Graph terminology



- vertices, edges, vertex degree, connected graph, tree
    - path, cycle, clique, spanning tree, spanning forest
  - Graph representations
    - array of edges
    - adjacency matrix
    - adjacency lists
  - Graph traversal
    - depth-first search (DFS)
    - breadth-first search (BFS)
    - cycle check, connected components
    - Hamiltonian paths/circuits, Euler paths/circuits
  - Suggested reading (Sedgewick):
    - graph representations ... Ch. 17.1-17.5
    - Hamiltonian/Euler paths ... Ch. 17.7
    - graph search ... Ch. 18.1-18.3, 18.7
-