

Week 8: Search Tree Algorithms

Tree Review

1/72

Binary search trees ...

- data structures designed for $O(\log n)$ search
- consist of nodes containing item (incl. key) and two links
- can be viewed as recursive data structure (subtrees)
- have overall ordering ($\text{data}(\text{Left}) < \text{root} < \text{data}(\text{Right})$)
- insert new nodes as leaves (or as root), delete from anywhere
- have structure determined by insertion order (*worst*: $O(n)$)
- operations: insert, delete, search, ...


Balanced Search Trees

Balanced BSTs

3/72

Reminder ...

- Goal: build binary search trees which have
 - minimum height \Rightarrow minimum worst case search cost
- Best balance you can achieve for tree with N nodes:
 - tree height of $\log_2 N \Rightarrow$ worst case search $O(\log N)$



Three *strategies* to improving worst case search in BSTs:

- *randomise* — reduce chance of worst-case scenario occurring
- *amortise* — do more work at insertion to make search faster
- *optimise* — implement all operations with performance bounds

Randomised BST Insertion

4/72

Effects of order of insertion on BST shape:

- best case (for at-leaf insertion): keys inserted in pre-order (median key first, then median of lower half, median of upper half, etc.)
- worst case: keys inserted in ascending/descending order
- average case: keys inserted in *random* order $\Rightarrow O(\log_2 n)$

Tree ADT has no control over order that keys are supplied.

Can the algorithm itself introduce some *randomness*?

In the hope that this randomness helps to balance the tree ...

... Randomised BST Insertion

5/72

How can a computer pick a number at random?

- it cannot

Software can only produce *pseudo random numbers*.

- a pseudo random number may appear unpredictable
 - but is actually predictable
- \Rightarrow implementation may deviate from expected theoretical behaviour
 - more on this in week 10

... Randomised BST Insertion

6/72

- Pseudo random numbers in C:

```
rand() // generates random numbers in the range 0 .. RAND_MAX
```

where the constant `RAND_MAX` is defined in `stdlib.h`
(depends on the computer: on the CSE network, `RAND_MAX` = 2147483647)

To convert the return value of `rand()` to a number between 0 .. RANGE

- compute the remainder after division by `RANGE+1`

... Randomised BST Insertion

7/72

Approach: normally do leaf insert, randomly do root insert.

```
insertRandom(tree,item)
|   Input tree, item
|   Output tree with item randomly inserted
|
|   if tree is empty then
|       return new node containing item
|   end if
|   // p/q chance of doing root insert
```

```

if random number mod  $q < p$  then
    return insertAtRoot(tree,item)
else
    return insertAtLeaf(tree,item)
end if

```

E.g. 30% chance \Rightarrow choose $p=3, q=10$

... Randomised BST Insertion

8/72

Cost analysis:

- similar to cost for inserting keys in random order: $O(\log_2 n)$
- does not rely on keys being supplied in random order

Approach can also be applied to deletion:

- standard method promotes inorder successor to root
- for the randomised method ...
 - promote inorder successor from right subtree, OR
 - promote inorder predecessor from left subtree

Rebalancing Trees

9/72

An approach to balanced trees:

- insert into leaves as for simple BST
- periodically, rebalance the tree

Question: how frequently/when/how to rebalance?

NewTreeInsert(tree,item):

```

Input  tree, item
Output tree with item randomly inserted

t=insertAtLeaf(tree,item)
if #nodes(t) mod  $k = 0$  then
    t=rebalance(t)
end if
return t

```

E.g. rebalance after every 20 insertions \Rightarrow choose $k=20$

Note: To do this efficiently we would need to change tree data structure and basic operations:

```


typedef struct Node {
    int data;
    int nnodes; // #nodes in my tree
    Tree left, right; // subtrees
} Node;

```

... Rebalancing Trees

10/72

How to rebalance a BST? Move median item to root.



... Rebalancing Trees

11/72

Implementation of rebalance:

```

rebalance(t):
    Input  tree t with n nodes
    Output t rebalanced

    if n ≥ 3 then
        t=partition(t, ⌊n/2⌋) // put node with median key at root
        left(t)=rebalance(left(t)) // then rebalance each subtree
        right(t)=rebalance(right(t))
    end if
    return t


```

... Rebalancing Trees

12/72

New operation on trees:

- **partition(tree,i)**: re-arrange tree so that element with index i becomes root




For tree with N nodes, indices are $0 \dots N-1$

... Rebalancing Trees

13/72

Partition: moves i^{th} node to root



... Rebalancing Trees

14/72

Implementation of partition operation:

```
partition(tree, i):
    Input tree with n nodes, index i
    Output tree with item #i moved to the root


    m = #nodes(left(tree))
    if i < m then
        left(tree) = partition(left(tree), i)
        tree = rotateRight(tree)
    else if i > m then
        right(tree) = partition(right(tree), i-m-1)
        tree = rotateLeft(tree)
    end if
    return tree
```

Note: size(tree) = n, size(left(tree)) = m, size(right(tree)) = n-m-1 (why -1?)


Exercise #1: Partition

15/72

Consider the tree t:



Show the result of partition(t, 3)



... Rebalancing Trees

17/72

Analysis of rebalancing: visits every node $\Rightarrow O(N)$

Cost means not feasible to rebalance after each insertion.

When to rebalance? ... Some possibilities:

- after every k insertions
- whenever "imbalance" exceeds threshold

Either way, we tolerate worse search performance for periods of time.

Does it solve the problem? ... Not completely \Rightarrow Solution: real balanced trees (later)

Splay Trees

19/72

A kind of "self-balancing" tree ...

Splay tree insertion modifies insertion-at-root method:

- by considering parent-child-grandchild (three level analysis)
- by performing double-rotations based on p-c-g orientation

The idea: appropriate double-rotations improve tree balance.

... Splay Trees

20/72

Splay tree implementations also do *rotation-in-search*:

- by performing double-rotations also when searching

The idea: provides similar effect to periodic rebalance.


\Rightarrow improves balance but makes search more expensive

... Splay Trees

21/72

Cases for splay tree double-rotations:


- case 1: grandchild is left-child of left-child \Rightarrow double right rotation from top
- case 2: grandchild is right-child of left-child
- case 3: grandchild is left-child of right-child
- case 4: grandchild is right-child of right-child \Rightarrow double left rotation from top



... Splay Trees

22/72

Double-rotation case for left-child of left-child ("zig-zig"):




Note: both rotations at the root (unlike insertion-at-root)

... Splay Trees

23/72

Double-rotation case for right-child of left-child ("zig-zag"):



Note: rotate subtree first (like insertion-at-root)

... Splay Trees

24/72

Algorithm for splay tree insertion:



```
insertSplay(tree,item):
    Input tree, item
    Output tree with item splay-inserted

    if tree is empty then return new node containing item
    else if item=data(tree) then return tree
    else if item<data(tree) then
        if left(tree) is empty then
            left(tree)=new node containing item
        else if item<data(left(tree)) then
            // Case 1: left-child of left-child "zig-zig"
            left(left(tree))=insertSplay(left(left(tree)),item)
            tree=rotateRight(tree)
        else if item>data(left(tree)) then
            // Case 2: right-child of left-child "zig-zag"
            right(left(tree))=insertSplay(right(left(tree)),item)
            left(tree)=rotateLeft(left(tree))
        end if
        return rotateRight(tree)
    else // item>data(tree)
        if right(tree) is empty then
            right(tree)=new node containing item
        else if item<data(right(tree)) then
            // Case 3: left-child of right-child "zag-zig"
            left(right(tree))=insertSplay(left(right(tree)),item)
            right(tree)=rotateRight(right(tree))
        else if item>data(right(tree)) then
            // Case 4: right-child of right-child "zag-zag"
            right(right(tree))=insertSplay(right(right(tree)),item)
            tree=rotateLeft(tree)
        end if
        return rotateLeft(tree)
    end if
```

Exercise #2: Splay Trees

25/72

Insert 18 into this splay tree:



... Splay Trees

27/72

Searching in splay trees:


```
searchSplay(tree,item):
|   Input  tree, item
|   Output address of item if found in tree
|           NULL otherwise
|
|   if tree=NULL then
|       return NULL
|   else
|       tree=splay(tree,item)
|       if data(tree)=item then
|           return tree
|       else
|           return NULL
|       end if
|   end if
```

where `splay()` is similar to `insertSplay()`,
except that it doesn't add a node ... simply moves `item` to root if found, or nearest node if not found


Exercise #3: Splay Trees

28/72

If we search for 22 in the splay tree



... how does this affect the tree?



... Splay Trees

30/72

Why take into account both child and grandchild?

- moves accessed node to the root
- *moves every ancestor of accessed node roughly halfway to the root*

⇒ better amortized cost than insert-at-root

... Splay Trees

31/72

Analysis of splay tree performance:

- assume that we "splay" for both insert and search
- consider: m insert+search operations, n nodes
- *Theorem.* Total number of comparisons: average $O((n+m) \cdot \log(n+m))$

Gives good overall (amortized) cost.

- insert cost not significantly different to insert-at-root
- search cost increases, but ...
 - improves balance on each search
 - moves frequently accessed nodes closer to root

But ... still has worst-case search cost $O(n)$

Real Balanced Trees

Better Balanced Binary Search Trees

33/72

So far, we have seen ...

- randomised trees ... make poor performance unlikely
- occasional rebalance ... fix balance periodically
- splay trees ... reasonable amortized performance
- but both types still have $O(n)$ worst case

Ideally, we want both average/worst case to be $O(\log n)$

- AVL trees ... fix imbalances as soon as they occur
- 2-3-4 trees ... use varying-sized nodes to assist balance
- red-black trees ... isomorphic to 2-3-4, but binary nodes

AVL Trees

AVL Trees

35/72

Invented by Georgy Adelson-Velsky and Evgenii Landis

Approach:

- insertion (at leaves) may cause imbalance
- repair balance as soon as we notice imbalance
- repairs done locally, not by overall tree restructure

A tree is unbalanced when: $\text{abs}(\text{height}(\text{left}) - \text{height}(\text{right})) > 1$

This can be repaired by at most two rotations:

- if left subtree too deep ...
 - if data inserted in left-right grandchild \Rightarrow left-rotate left subtree
 - rotate right
- if right subtree too deep ...
 - if data inserted in right-left grandchild \Rightarrow right-rotate right subtree
 - rotate left

Problem: determining height/depth of subtrees may be expensive.

... AVL Trees

36/72

Implementation of AVL insertion


```
insertAVL(tree, item):  
| Input   tree, item  
| Output tree with item AVL-inserted
```

```
if tree is empty then  
    return new node containing item  
else if item=data(tree) then  
    return tree  
else  
    if item<data(tree) then  
        left(tree)=insertAVL(left(tree),item)  
    else if item>data(tree) then  
        right(tree)=insertAVL(right(tree),item)  
    end if  
    if height(left(tree))-height(right(tree)) > 1 then  
        if item>data(left(tree)) then  
            left(tree)=rotateLeft(left(tree))  
        end if  
        tree=rotateRight(tree)  
    else if height(right(tree))-height(left(tree)) > 1 then  
        if item<data(right(tree)) then  
            right(tree)=rotateRight(right(tree))  
        end if  
        tree=rotateLeft(tree)  
    end if  
    return tree  
end if
```

Exercise #4: AVL Trees

37/72

Insert 27 into the AVL tree



What would happen if you now insert 28?

You may like the animation at www.cs.usfca.edu/~galles/visualization/AVLtree.html


... AVL Trees

39/72

Analysis of AVL trees:

- trees are *height*-balanced; subtree depths differ by ± 1

- average/worst-case search performance of $O(\log n)$
- *require* extra data to be stored in each node ("height")
- may not be *weight*-balanced; subtree sizes may differ




2-3-4 Trees

2-3-4 Trees

41/72

2-3-4 trees have three kinds of nodes

- 2-nodes, with two children (same as normal BSTs)
- 3-nodes, two values and three children
- 4-nodes, three values and four children



... 2-3-4 Trees

42/72

2-3-4 trees are ordered similarly to BSTs



In a *balanced* 2-3-4 tree:

- all leaves are at same distance from the root

2-3-4 trees grow "upwards" by splitting 4-nodes.

... 2-3-4 Trees

43/72

Possible 2-3-4 tree data structure:

```
typedef struct node {
    int         order;      // 2, 3 or 4
    int         data[3];    // items in node
    struct node *child[4];  // links to subtrees
} node;
```

... 2-3-4 Trees

44/72

Searching in 2-3-4 trees:

Search(tree,item):

Input tree, item

Output address of item if found in 2-3-4 tree
NULL otherwise

if tree is empty **then**
 return NULL

else

 i=0
 while i<tree.order-1 and item>tree.data[i] **do**
 i=i+1 // find relevant slot in data[]
 end while

if item=tree.data[i] **then** // item found
 return address of tree.data[i]

else // keep looking in relevant subtree
 return Search(tree.child[i],item)

end if

end if

... 2-3-4 Trees

45/72

2-3-4 tree searching cost analysis:

- as for other trees, worst case determined by height h
- 2-3-4 trees are always balanced \Rightarrow height is $O(\log n)$
- worst case for height: all nodes are 2-nodes
 same case as for balanced BSTs, i.e. $h \approx \log_2 n$
- best case for height: all nodes are 4-nodes
 balanced tree with branching factor 4, i.e. $h \approx \log_4 n$

Insertion into 2-3-4 Trees

46/72

Starting with the root node:

repeat

- if current node is full (i.e. contains 3 items)
 - split into two 2-nodes
 - promote middle element to parent
 - if no parent \Rightarrow middle element becomes the new root 2-node


- go back to parent node
- if current node is a leaf
 - insert Item in this node, order++
- if current node is not a leaf
 - go to child where Item belongs

until Item inserted

... Insertion into 2-3-4 Trees

47/72

Insertion into a 2-node or 3-node:




Insertion into a 4-node (requires a split):



... Insertion into 2-3-4 Trees

48/72


Splitting the root:



... Insertion into 2-3-4 Trees

49/72



Building a 2-3-4 tree ... 7 insertions:



Exercise #5: Insertion into 2-3-4 Tree

50/72

Show what happens when D, S, F, U are inserted into this tree:



... Insertion into 2-3-4 Trees

52/72

Insertion algorithm:

```
insert(tree,item):
  Input  2-3-4 tree, item
  Output tree with item inserted

  node=root(tree), parent=NULL
  repeat
    if node.order=4 then
      promote = node.data[1]    // middle value
      nodeL   = new node containing node.data[0]
      nodeR   = new node containing node.data[2]
      if parent=NULL then
        make new 2-node root with promote,nodeL,nodeR
      else
        insert promote,nodeL,nodeR into parent
        increment parent.order
      end if
      node=parent
    end if
    if node is a leaf then
      insert item into node
```



```

    increment node.order
  else
    parent=node
    if item<node.data[0] then
      node=node.child[0]
    else if item<node.data[1] then
      node=node.child[1]
    else
      node=node.child[2]
    end if
  end if
end if
until item inserted

```

... Insertion into 2-3-4 Trees

53/72

Variations on 2-3-4 trees ...

Variation #1: why stop at 4? why not 2-3-4-5 trees? or M -way trees?

- allow nodes to hold up to $M-1$ items, and at least $M/2$
- if each node is a disk-page, then we have a *B-tree* (databases)
- for B-trees, depending on Item size, $M > 100/200/400$

Variation #2: don't have "variable-sized" nodes

- use standard BST nodes, augmented with one extra piece of data
- implement similar strategy as 2-3-4 trees \rightarrow red-black trees.

Red-Black Trees

Red-Black Trees

55/72

Red-black trees are a representation of 2-3-4 trees using BST nodes.

- each node needs one extra value to encode link type
- but we no longer have to deal with different kinds of nodes

Link types:

- *red* links ... combine nodes to represent 3- and 4-nodes
- *black* links ... analogous to "ordinary" BST links (child links)

Advantages:

- standard BST search procedure works unmodified
- get benefits of 2-3-4 tree self-balancing (although deeper)

Red-Black Trees

56/72

Definition of a *red-black tree*

- a BST in which each node is marked red or black
- no two red nodes appear consecutively on any path
- a red node corresponds to a 2-3-4 sibling of its parent
- a black node corresponds to a 2-3-4 child of its parent

Balanced red-black tree

- all paths from root to leaf have same number of black nodes


Insertion algorithm: avoids worst case $O(n)$ behaviour

Search algorithm: standard BST search

... Red-Black Trees

57/72

Representing 4-nodes in red-black trees:




Some texts colour the links rather than the nodes.

... Red-Black Trees

58/72


Representing 3-nodes in red-black trees (two possibilities):




... Red-Black Trees

59/72

Equivalent trees (one 2-3-4, one red-black):



- black = parent node is a "real" parent
- red = parent node is a 2-3-4 neighbour



... Red-Black Trees

60/72

Red-black tree implementation:

```
typedef enum {RED,BLACK} Colr;
typedef struct node *RBTree;
typedef struct node {
    int    data;    // actual data
    Colr   colour;  // relationship to parent
    RBTree left;    // left subtree
    RBTree right;   // right subtree
} node;

#define colour(tree) ((tree)->colour)
#define isRed(tree)  ((tree) != NULL && (tree)->colour == RED)
```

RED = node is part of the same 2-3-4 node as its parent (sibling)

BLACK = node is a child of the 2-3-4 node containing the parent

... Red-Black Trees

61/72

New nodes are always red:

```
RBTree newNode(Item it) {
    RBTree new = malloc(sizeof(Node));
    assert(new != NULL);
    data(new) = it;
    colour(new) = RED;
    left(new) = right(new) = NULL;
    return new;
}
```

... Red-Black Trees

62/72

Node.colour allows us to distinguish links

... Red-Black Trees

63/72

Search method is standard BST search:

```
SearchRedBlack(tree,item):
    Input  tree, item
    Output true if item found in red-black tree
           false otherwise

    if tree is empty then
        return false
    else if item < data(tree) then
        return SearchRedBlack(left(tree),item)
    else if item > data(tree) then
        return SearchRedBlack(right(tree),item)
    else // found
        return true
    end if
```

Red-Black Tree Insertion

64/72

Insertion is more complex than for standard BSTs

- need to recall direction of last branch (L or R)
- need to recall whether parent link is red or black
- splitting/promoting implemented by rotateLeft/rotateRight
- several cases to consider depending on colour/direction combinations

... Red-Black Tree Insertion

65/72

High-level description of insertion algorithm:

```
insertRB(tree,item,inRight):
    Input  tree, item, inRight indicating direction of last branch
    Output tree with it inserted

    if tree is empty then
        return newNode(item)
    else if item = data(tree) then
```

```

    return tree
end if
if left(tree) and right(tree) both are RED then
    split 4-node in a red-black tree
end if
recursive insert a la BST, re-arrange links/colours after insert
return modified tree

```

```

insertRedBlack(tree,item):
    Input  red-black tree, item
    Output tree with item inserted


    tree=insertRB(tree,item,false)
    colour(tree)=BLACK
    return tree

```

... Red-Black Tree Insertion

66/72

Splitting a 4-node, in a red-black tree:



Algorithm:

```

colour(currentTree)=RED
colour(left(currentTree))=BLACK
colour(right(currentTree))=BLACK

```

... Red-Black Tree Insertion

67/72

Simple recursive insert (a la BST):



Algorithm:

```

if item<data(tree) then
    left(tree)=insertRB(left(tree),item,false)
    re-arrange links/colours after insert
else
    // item larger than data in root
    right(tree)=insertRB(right(tree),item,true)

```

```

    re-arrange links/colours after insert
end if

```

Not affected by colour of tree node.

... Red-Black Tree Insertion

68/72

Re-arrange links/colours after insert:

Step 1 — "normalise" direction of successive red links



Algorithm:

```

if inRight and both currentTree and left(currentTree) are red then
    currentTree=rotateRight(currentTree)
end if

```

Symmetrically,

- if not inRight and both currentTree and right(currentTree) are red
⇒ left rotate currentTree

... Red-Black Tree Insertion

69/72

Re-arrange links/colours after insert:

Step 2 — two successive red links = newly-created 4-node



Algorithm:

```

if left(currentTree) and left(left(currentTree)) are red then
    currentTree=rotateRight(currentTree)
    colour(currentTree)=BLACK
    colour(right(currentTree))=RED
end if


```

Symmetrically,

- if both right(currentTree) and right(right(currentTree)) are red
⇒ left rotate currentTree, then re-colour currentTree and left(currentTree)

Example of insertion, starting from empty tree:

22, 12, 8, 15, 11, 19, 43, 44, 45, 42, 41, 40, 39



Red-black Tree Performance

71/72

Cost analysis for red-black trees:

- tree is well-balanced; worst case search is $O(\log_2 n)$
- insertion affects nodes down one path; max #rotations is $2 \cdot h$
(where h is the height of the tree)

Only disadvantage is complexity of insertion/deletion code.

Note: red-black trees were popularised by Sedgewick.

Summary

72/72

- Randomised at-leaf/at-root insertion
- Tree operations
 - tree partition
 - joining trees
- Self-adjusting trees
 - Splay trees
 - AVL trees
 - 2-3-4 trees
 - Red-black trees
- Suggested reading:
 - Sedgewick, Ch. 12.9
 - Sedgewick, Ch. 13.1-13.4